

Lecture # 39

- Instruction Set Architecture
- CISC vs. RISC

Electronic Discrete Variable Automatic Computer (EDVAC)

- ENIAC' s programming system was external
 - Sequences of instructions were executed independently of the results of the calculation
 - Human intervention required to take instructions “out of order”
- Eckert, Mauchly, John von Neumann and others designed EDVAC (1944) to solve this problem
 - Solution was the *stored program computer*
 - ® “*program can be manipulated as data*”
- *First Draft of a report on EDVAC* was published in 1945, but just had von Neumann' s signature!
 - In 1973 the court of Minneapolis attributed the honor of *inventing the computer* to John Atanasoff

And then there was IBM 701

IBM 701 -- 30 machines were sold in 1953-54

IBM 650 -- a cheaper, drum based machine,
more than 120 were sold in 1954
and there were orders for 750 more!
- eventually sold about 2000 of them

Users stopped building their own machines.

Why was IBM late getting into computer technology?

IBM was making too much money!

Even without computers, IBM revenues were doubling every 4 to 5 years in 40's and 50's.

Into the 60's...:

Compatibility Problem at IBM

By early 60's, IBM had 4 incompatible lines of computers!

701	□	7094
650	□	7074
702	□	7080
1401	□	7010

Each system had its own

- Instruction set
- Peripherals: magnetic tapes, drums and disks
- Programming tools: assemblers, compilers, libraries,...
- market niche: business, scientific, etc....

IBM 360 : Design Premises

Amdahl, Blaauw and Brooks, 1964

<http://www.research.ibm.com/journal/rd/441/amdahl.pdf>

- Breaks the link between programmer and hardware
- Upward and downward, machine-language compatibility across a family of machines
- General purpose machine organization, general I/O interfaces, storage > 32K
- Easier to use (answers-per-month vs. bits-per-second)
- Machine must be capable of *supervising itself* without manual intervention → OS/360 (simple OS' s in IBM 700/7000)
- Built-in *hardware fault checking* and locating aids to reduce down time


... the use of the “ISA” as a compatibility layer
was a \$175 billion project (2011 dollars)

The Amdahl .. from Amdahl's Law.

The Brooks .. from The Mythical Man-Month.

The Stored Program Computer

- The program is *data*
 - It is a series of bits
 - It lives in memory
 - A series of discrete “instructions”
- The program counter (PC) control execution
 - It points to the current instruction
 - Advances through the program



Instruction Memory			
[80000180]	0001d821	addu \$27, \$0, \$1	
[80000184]	3c019000	lui \$1, -28672	
[80000188]	ac220200	sw \$2, 512(\$1)	
[8000018c]	3c019000	lui \$1, -28672	
[80000190]	ac240204	sw \$4, 516(\$1)	
[80000194]	401a6800	mfc0 \$26, \$13	
[80000198]	001a2082	srl \$4, \$26, 2	
[8000019c]	3084001f	andi \$4, \$4, 31	
[800001a0]	34020004	ori \$2, \$0, 4	
[800001a4]	3c049000	lui \$4, -28672 [__m1_]	
[800001a8]	0000000c	syscall	
[800001ac]	34020001	ori \$2, \$0, 1	
[800001b0]	001a2082	srl \$4, \$26, 2	
[800001b4]	3084001f	andi \$4, \$4, 31	

Data Memory			
[7ffffe60]	74736574	73612e32	t e s t 2 . a s
[7ffffe70]	5f524553	54584554	S E R _ T E X T
[7ffffe80]	78303d47	3a364631	G = 0 x 1 F 6 :
[7ffffe90]	5f444e41	45444f4d	A N D _ M O D E
[7ffffea0]	70410033	5f656c70	3 A p p l e _ 1
[7ffffeb0]	656b636f	65525f74	o c k e t _ R e
[7ffffec0]	616c2f70	68636e75	p / l a u n c h

The Instruction Set Architecture (ISA)

- The ISA is the set of instructions a computer can execute
- All programs are combinations of these instructions
- It is an abstraction that programmers (and compilers) use to express computations
 - The ISA defines a set of operations, their semantics, and rules for their use.
 - The software agrees to follow these rules.
- The hardware can implement those rules **IN ANY WAY IT CHOOSES!**
 - Directly in hardware
 - Via a software layer (i.e., a virtual machine)
 - Via a trained monkey with a pen and paper
 - Via a software simulator (like SPIM)

RISC vs CISC

In the Beginning...

- 1964 -- The first ISA appears on the IBM System 360
- In the “good” old days
 - Initially, the focus was on usability by humans.
 - Lots of “user-friendly” instructions (remember the x86 addressing modes).
 - Memory was expensive, so code-density mattered.
 - Many processors were *microcoded* -- each instruction actually triggered the execution of a builtin function in the CPU. Simple hardware to execute complex instructions (but CPIs are very, very high)
- ...SO...
 - Many, many different instructions, lots of bells and whistles
 - Variable-length instruction encoding to save space.
- ... their success had some downsides...
 - ISAs evolved organically.
 - They got messier, and more complex.

Things Changed

- In the modern era
 - Compilers write code, not humans.
 - Memory is cheap. Code density is unimportant.
 - Low CPI should be possible, but only for simple instructions
 - We learned a lot about how to design ISAs, how to let them evolve gracefully, etc.
- So, architects started with with a clean slate...

Reduced Instruction Set Computing (RISC)

- Simple, regular ISAs, mean simple CPUs, and simple CPUs can go fast.
 - Fast clocks.
 - Low CPI.
 - Simple ISAs will also mean more instruction (increasing IC), but the benefits should outweigh this.
- Compiler-friendly, not user-friendly.
 - Simple, regular ISAs, will be easy for compilers to use
 - A few, simple, flexible, fast operations that compiler can combine easily.
 - Separate memory access and data manipulation
 - Instructions access memory *or* manipulate register values. Not both.
 - “Load-store architectures” (like MIPS)

RISC Characteristics of MIPS

- All instructions have
 - ≤ 1 arithmetic op
 - ≤ 1 memory access
 - ≤ 2 register reads
 - ≤ 1 register write
 - ≤ 1 branch
 - It needs a small, fixed amount of hardware.
- Instructions operate on memory *or* registers *not both*
 - “Load/Store Architecture”
- Decoding is easy
 - Uniform opcode location
 - Uniform register location
 - Always 4 bytes -> the location of the next PC is to know.
- Uniform execution algorithm
 - Fetch
 - Decode
 - Execute
 - Memory
 - Write Back
- Compiling is easy
 - No complex instructions to reason about
 - No special registers
- The HW is simple
 - A skilled undergrad can build one in 10 weeks.
 - 33 instructions can run complex programs.

CISC: x86

- x86 is the prime example of CISC (there were many others long ago)
 - Many, many instruction formats. Variable length.
 - Many complex rules about which register can be used when, and which addressing modes are valid where.
 - Very complex instructions
 - Combined memory/arithmetic.
 - Special-purpose registers.
 - Many, many instructions.
- Implementing x86 correctly is almost intractable

x86 ISA Caveats

- x86 is a poorly-designed ISA
 - It breaks almost every rule of good ISA design.
 - There is nothing “regular” or predictable about its syntax.
 - We don’ t have time to learn how to write x86 with any kind of thoroughness.
- It is the most widely used ISA in the world today.
 - It is the ISA you are most likely to see in the “real world”
 - So it’ s useful to study.
- Intel and AMD have managed to engineer (at considerable cost) their CPUs so that this ugliness has relatively little impact on their processors’ performance (more on this later)